

Steps for Introducing a System Call

- Ensure that you are in the linux-xx.xx.y directory
- Create your own directory
 - mkdir info
 - cd info/
 - **Create processInfo.h**
 - asmlinkage long sys_listProcessInfo(void);
 - Create processInfo.c
 - Write Makefile and change top level Makefile
 - obj-y:=listProcessInfo.o
 - core -y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/ info/
 - Alter the /arch/x86/entry/syscalls/syscall_64.tbl
 - 548 common processInfo sys_processInfo
 - Alter /include/linux/syscalls.h
 - Add asmlinkage long sys_hello(void)
 - **Compile**
 - sudo make && sudo make modules_install && sudo make install
 - **Test**

Process

- Definition: **execution environment with restricted rights**
 - Address Space with One or More Threads
 - Owns memory (mapped pages)
 - Owns file descriptors, file system context, ...
 - Encapsulates one or more threads sharing process resources
- Application program executes as a process
 - Complex applications can fork/exec child processes [later]
- Why processes?
 - Protected from each other. OS Protected from them.
 - Execute concurrently [trade-offs with threads? later]
 - Basic unit OS deals with

Process

- Definition: **execution environment with restricted rights**

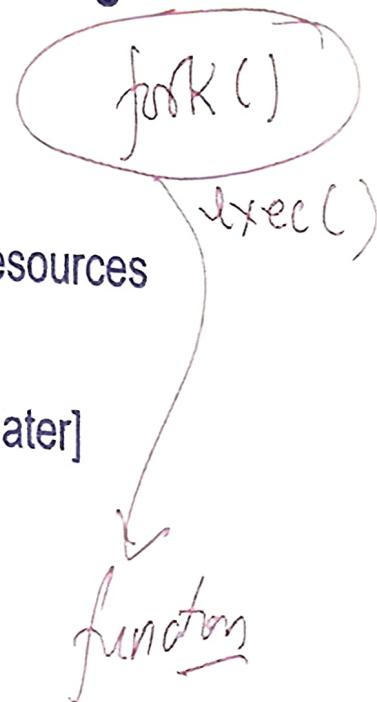
- Address Space with One or More Threads
- Owns memory (mapped pages)
- Owns file descriptors, file system context, ...
- Encapsulates one or more threads sharing process resources

- Application program executes as a process

- Complex applications can fork/exec child processes [later]

- Why processes?

- Protected from each other. OS Protected from them.
- Execute concurrently [trade-offs with threads? later]
- Basic unit OS deals with



pid = fork()

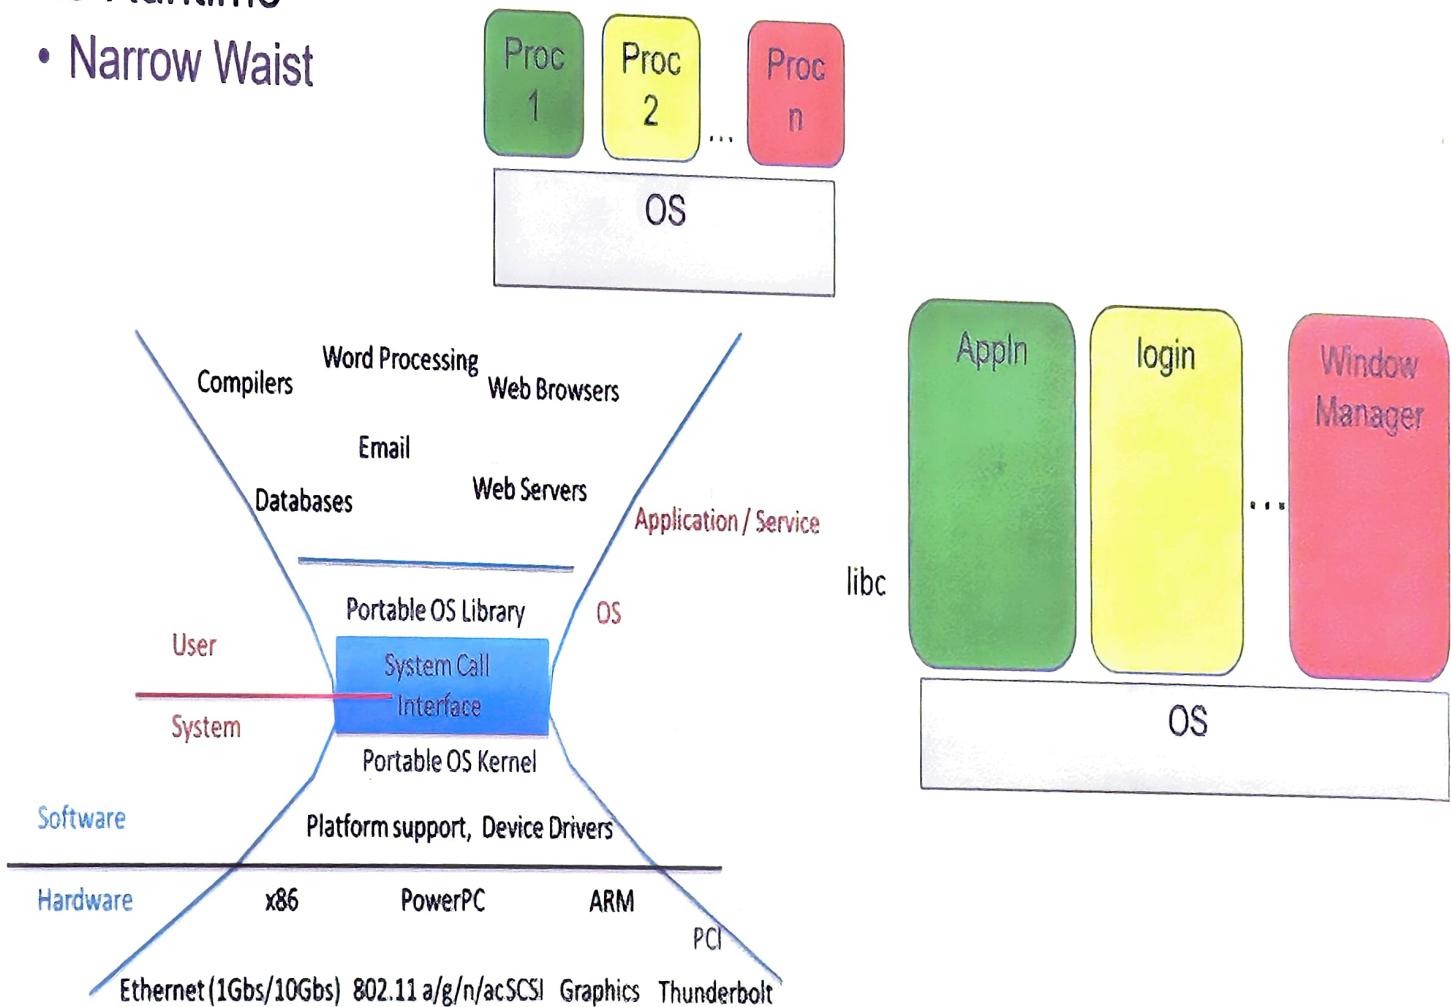
if (pid == 0) do something

else

exec()

OS Runtime Library

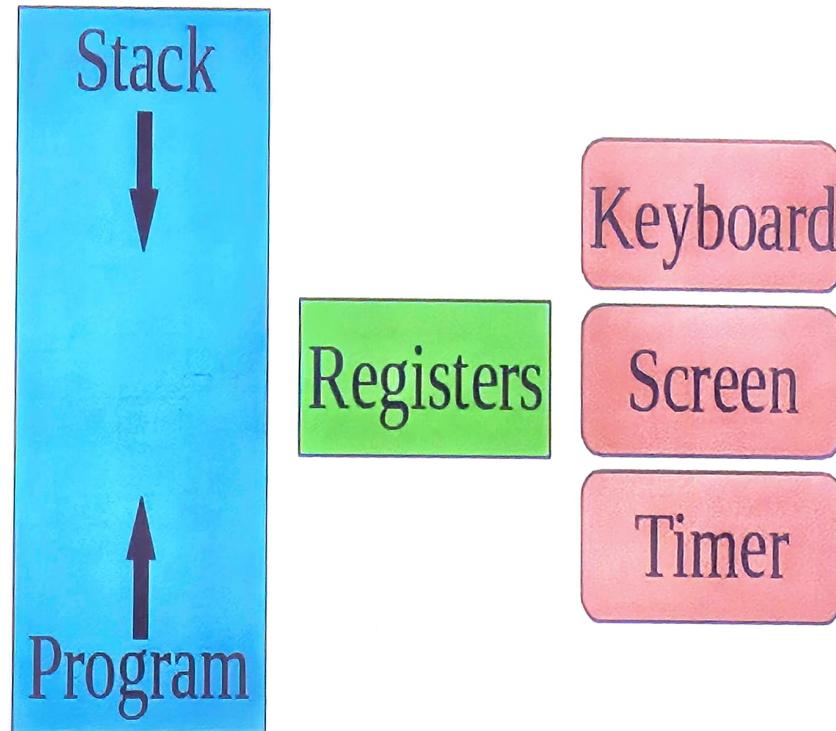
- The Runtime
 - Narrow Waist



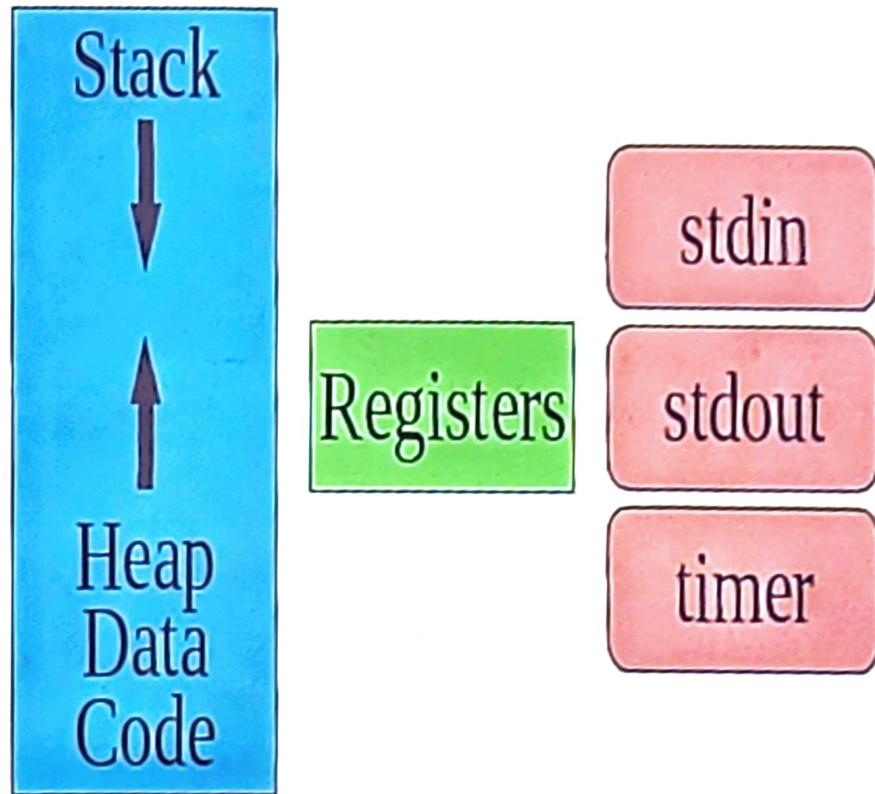
Process Concept

- An operating system executes a variety of programs:
 - Batch system – jobs
 - Time-shared systems – user programs or tasks
- Textbook uses the terms *job* and *process* almost interchangeably
- Process – a program in execution; process execution must progress in sequential fashion
- A process includes:
 - program counter
 - stack
 - data section

The Bare Machine



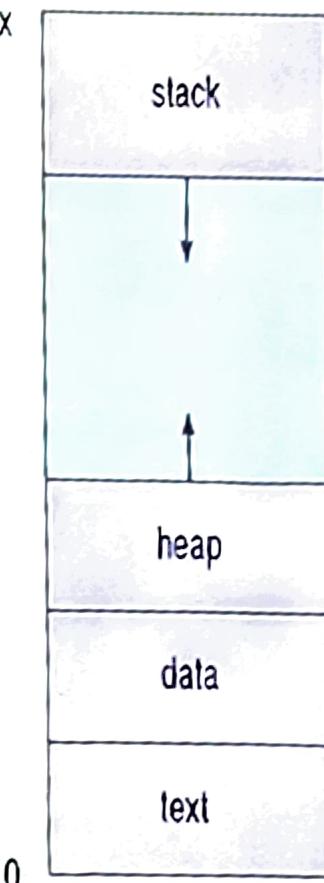
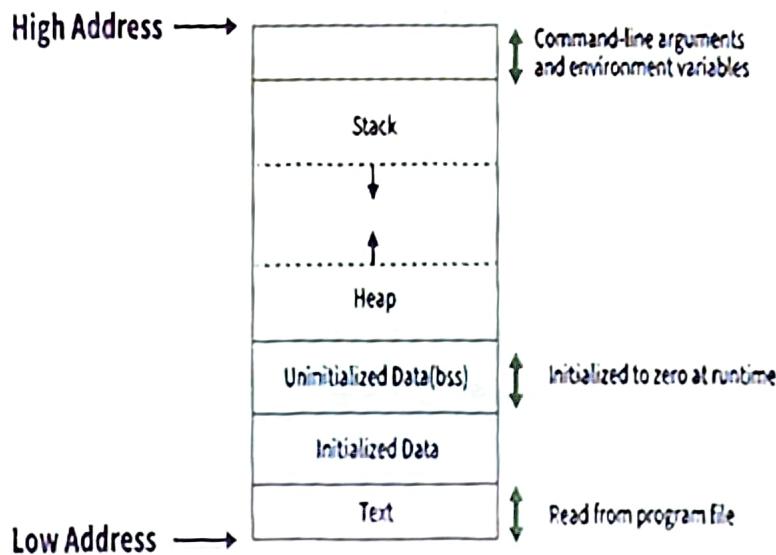
The Bare Machine



Process in Memory

- Kernel builds stack for new process

- Transfers argv[] and envp[] to top of new process^{max}
- Hand-crafts stack frame for __main()
- Sets registers
 - Stack pointer (to top frame)
 - Program counter (to start of __main())



```
(base) kolin@mosaic:~$  
(base) kolin@mosaic:~$ vi test.c  
(base) kolin@mosaic:~$ gcc test.c -o t1  
(base) kolin@mosaic:~$ size t1  
text    data    bss    dec    hex filename  
1228     544      8    1780    6f4 t1  
(base) kolin@mosaic:~$ vi test.c  
(base) kolin@mosaic:~$ gcc test.c -o t1  
(base) kolin@mosaic:~$ size t1  
text    data    bss    dec    hex filename  
1228     544      8    1780    6f4 t1  
(base) kolin@mosaic:~$ vi test.c  
  
(base) kolin@mosaic:~$ gcc test.c -o t1  
(base) kolin@mosaic:~$ size t1  
text    data    bss    dec    hex filename  
1228     544      8    1780    6f4 t1  
(base) kolin@mosaic:~$
```



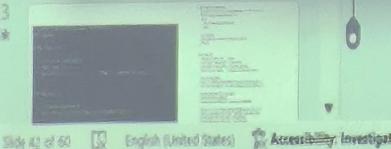
```
int g_a=0;
int main(){
int a;
    printf("%d\n",g_a+a);

    return 0;
}
```

3,6

All

43



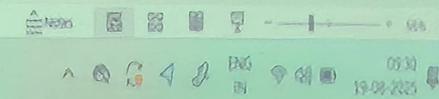
Slide 43 of 60

English (United States)

Accessibility, Investigations



Search



```
(base) kolin@mosaic:~$ size t1
text    data    bss    dec    hex filename
1228     544      8   1780    6f4 t1
(base) kolin@mosaic:~$ vi test.c
(base) kolin@mosaic:~$ conda deactivate
kolin@mosaic:~$
kolin@mosaic:~$
kolin@mosaic:~$ vi test.c
kolin@mosaic:~$ gcc test.c -o t1
kolin@mosaic:~$ size t1
text    data    bss    dec    hex filename
1228     544      8   1780    6f4 t1
kolin@mosaic:~$ vi test.c
kolin@mosaic:~$ gcc test.c -o t1
kolin@mosaic:~$ size t1
text    data    bss    dec    hex filename
1397     600      8   2005    7d5 t1
kolin@mosaic:~$
```

```
// uninitialized global variable
int ugvar;

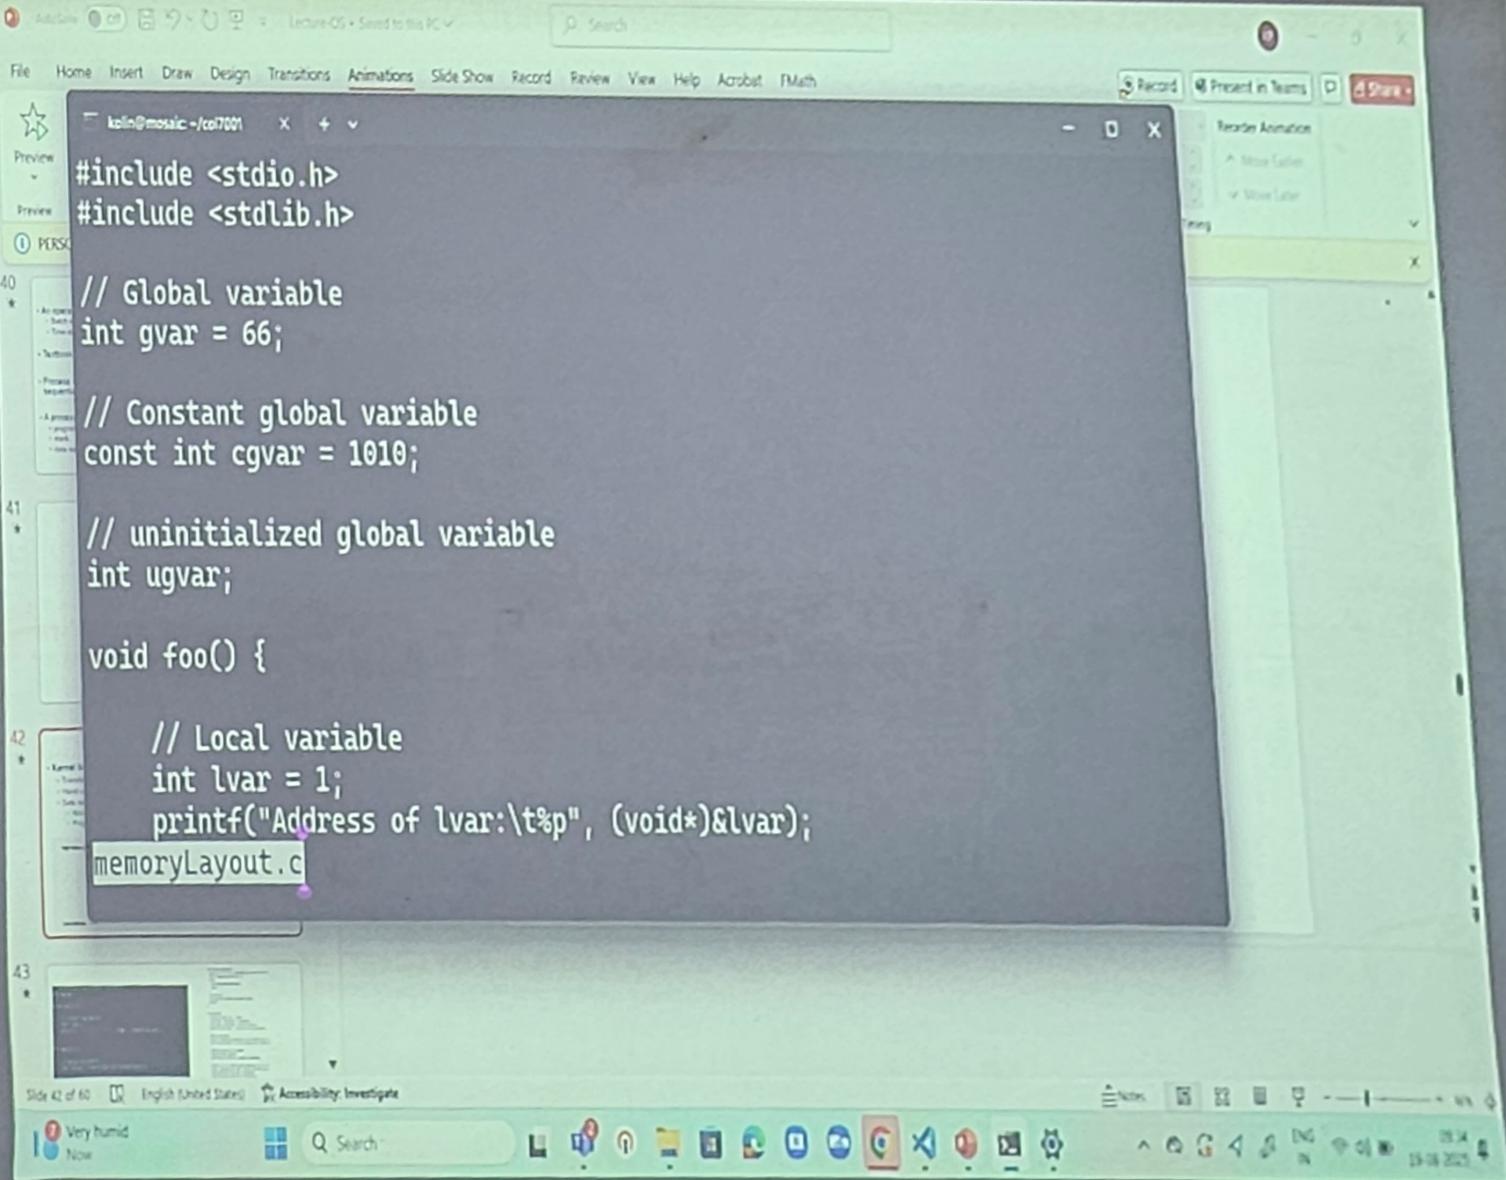
void foo() {

    // Local variable
    int lvar = 1;
    printf("Address of lvar:\t%p\n", (void*)&lvar);
}

int main() {

    // Heap variable
    int *hvar = (int*)malloc(sizeof(int));

    // Checking and comparing address of different
    // elements of program that should be stored in
    // different segments of the memory
    printf("Address of foo:\t%p\n", (void*)&foo);
    printf("Address of cgvar:\t%p\n", (void*)&cgvar);
    printf("Address of gvar:\t%p\n", (void*)&gvar);
    printf("Address of ugvar:\t%p\n", (void*)&ugvar);
    printf("Address of hvar:\t%p\n", (void*)hvar);
    foo();
    printf("\n\n");
    return 0;
}
```



Lecture-05 • Saved to this PC

File Home Insert Draw Design Transitions Animations Slide Show Record Review View Help Acrobat FMath

Record Present in Teams Share

Preview

PERSO

40

41

42

43

kolin@mosaic:~/col7001

| | text | data | bss | dec | hex | filename |
|--|------|------|-----|------|-----|----------|
| | 1381 | 600 | 8 | 1989 | 7c5 | t1 |

```
kolin@mosaic:~$  
kolin@mosaic:~$  
kolin@mosaic:~$  
kolin@mosaic:~$ ls  
accel  
anaconda3  
BaseNet  
col7001  
Desktop  
dna_r9.4.1_e8-FLO_FLG001-SQK_RAB204-4000.pod5  
Documents  
Downloads  
genomicAccelerator  
kolin@mosaic:~$ cd col7001/  
kolin@mosaic:~/col7001$ less memoryLayout.c  
kolin@mosaic:~/col7001$ gcc
```

Reorder Animation

Mouse Earlier

Mouse Later

Timing

Notes

Search

Autosave off

English (United States)

Accessibility: Investigate

Slide 42 of 60

NIFTY +0.06%

ENG IN 23-08-2025 09:35



Preview



PERSON

40

- An open file
- Read
- Write
- Text file
- Plain text
- Plain text with comments
- A process
- program
- work
- done

41

Downloads

42



Kernel

↳

Translators

↳

Host

↳

Sudo

↳

Programs

↳

kolin@mosaic:~/col7001

kolin@mosaic:~\$ ls

accel

anaconda3

BaseNet

col7001

Desktop

dna_r9.4.1_e8-FLO_FLG001-SQK_RAB204-4000.pod5

Documents

Downloads

genomicAccelerator

logs

misc

Music

Pictures

pptx

Public

Quantum

Sample.tsv

satmap

security

snap

t1

Templates

test.c

Untitled.ipynb

Videos

kolin@mosaic:~\$ cd col7001/

kolin@mosaic:~/col7001\$ less memoryLayout.c

kolin@mosaic:~/col7001\$ gcc -o memoryLayout memoryLayout.c

kolin@mosaic:~/col7001\$./memoryLayout^C

kolin@mosaic:~/col7001\$ size memoryLayout

text data bss dec hex filename

2097 628 12 2737 ab1 memoryLayout

kolin@mosaic:~/col7001\$

kolin@mosaic:~/col7001

Downloads

genomicAccelerator

0 PERSON

Sample.tsv

satmap

kolin@mosaic:~\$ cd col7001/

kolin@mosaic:~/col7001\$ less memoryLayout.c

kolin@mosaic:~/col7001\$ gcc -o memoryLayout memoryLayout.c

kolin@mosaic:~/col7001\$./memoryLayout^C

kolin@mosaic:~/col7001\$ size memoryLayout

| text | data | bss | dec | hex | filename |
|------|------|-----|-----|-----|----------|
|------|------|-----|-----|-----|----------|

| | | | | | |
|------|-----|----|------|-----|--------------|
| 2097 | 628 | 12 | 2737 | ab1 | memoryLayout |
|------|-----|----|------|-----|--------------|

kolin@mosaic:~/col7001\$./memoryLayout

Address of foo: 0x5555555551a9

Address of cgvar: 0x5555555556004

Address of gvar: 0x5555555558010

Address of ugvar: 0x5555555558018

Address of hvar: 0x55555555592a0

Address of lvar: 0x7fffffff294

kolin@mosaic:~/col7001\$

```

// uninitialized global variable
int ugvar;

void foo() {

    // Local variable
    int lvar = 1;
    printf("Address of lvar:\t%p", (void*)&lvar);
}

int main() {

    // Heap variable
    int *hvar = (int*)malloc(sizeof(int));

    // Checking and comparing address of different
    // elements of program that should be stored in
    // different segments of the memory
    printf("Address of foo:\t%p\n", (void*)&foo);
    printf("Address of cgvar:\t%p\n", (void*)&cgvar);
    printf("Address of gvar:\t%p\n", (void*)&gvar);
    printf("Address of ugvar:\t%p\n", (void*)&ugvar);
    printf("Address of hvar:\t%p\n", (void*)&hvar);
    foo();
    printf("\n\n");
    return 0;
}

```

```

static void print_maps(void) {
    puts("***** /proc/self/maps *****");
    FILE *f = fopen("/proc/self/maps", "r");
    if (!f) {
        perror("fopen(/proc/self/maps)");
        return;
    }
    char line[512];
    while (fgets(line, sizeof line, f)) { /*getline, sic! */
        close(f);
        int main(void) {
            int local_on_stack = 123; // stack
            void *heap1 = malloc(1234); // heap probe
            void *heap2 = malloc(1234); // another allocation
            void *brk_now = sbrk(0); // program break (heap end here)

            char eexec_path[4096];
            size_t i = 0;
            readlink("/proc/self/exe", eexec_path, sizeof eexec_path - 1);
            if (i > 0) { eexec_path[i] = '\0'; } else { strcpy(eexec_path, "unknown"); }

            printf("Executable %s\n", eexec_path);
            printf("FD %d\n", getpid());
            printf("Page size %d bytes\n", sysconf(_SC_PAGESIZE));

            puts("***** High-level regions (addresses) *****");
            printf("_executable_start: %p\n", (void*)A_executable_start);
            printf("end (end of .text): %p\n", (void*)A_end);
            printf("_edata [end of .data]: %p\n", (void*)A_end);
            printf("end (end of .bss): %p\n", (void*)B_end);

            puts("\n----- Probes (examples in these regions) -----");
            printf("text main : %p\n", (void*)A_main);
            printf("text foo : %p\n", (void*)A_foo);
            printf("data crt : %p -> (%p, %p)\n", (void*)A_crt, (void*)C_start, C_end);
            printf("data __init : %p (%p)\n", (void*)A_init, (void*)A_end);
            printf("bss _uninit : %p\n", (void*)B_uninit);
            printf("heap malloc(10) : %p, %p\n", (void*)B_heap1, B_heap2);
            printf("heap program brk: %p (%p)\n", (void*)B_brk, B_brk_now);
            printf("stack local var : %p (%p)\n", (void*)B_local_on_stack, (void*)B_local_on_stack);

            puts("\n----- Layout notes -----");
            puts("On most x86_64 Linux builds with PIE, the executable is position-independent and "
                 "the loader maps segments at randomized base addresses ([3,4]).");
            puts(" _executable_start < text or _text < data or _data < bss or end.");
            puts("Heap grows upward from the program break, stack grows downward from high addresses.");
            puts("For the authoritative mapping (including shared lib), see /proc/self/maps below.");
        }
        print_maps();
    }
    free(heap1);
    free(heap2);
    return 0;
}

```

```
static void print_maps(void) {
    puts("\n===== /proc/self/maps =====");
    FILE *f = fopen("/proc/self/maps", "r");
    if (!f) {
        perror("fopen(/proc/self/maps)");
        return;
    }
    char line[512];
    while (fgets(line, sizeof line, f)) fputs(line, stdout);
    fclose(f);
}

int main(void) {
    int local_on_stack = 123; // stack
    void *heap1 = malloc(1024); // heap probe
    void *heap2 = malloc(1024); // another allocation
    void *brk_now = sbrk(0); // program break (heap end hint)

    char exe_path[4096];
    ssize_t n = readlink("/proc/self/exe", exe_path, sizeof exe_path - 1);
    if (n >= 0) { exe_path[n] = '\0'; } else { strcpy(exe_path, "unknown"); }

    printf("Executable:%s\n", exe_path);
    printf("PID:%d\n", getpid());
    printf("Page size: %ld bytes\n", sysconf(_SC_PAGESIZE));
```

```
int local_on_stack = 123; // stack
void *heap1 = malloc(1024); // heap probe
void *heap2 = malloc(1024); // another allocation
void *brk_now = sbrk(0); // program break (heap end hint)

char exe_path[4096];
ssize_t n = readlink("/proc/self/exe", exe_path, sizeof(exe_path) - 1);
if (n >= 0) { exe_path[n] = '\0'; } else { strcpy(exe_path, "unknown"); }

printf("Executable: %s\n", exe_path);
printf("PID: %d\n", getpid());
printf("Page size: %ld bytes\n", sysconf(_SC_PAGESIZE));

puts("\n===== High-level regions (addresses) =====");
printf("_executable_start: %p\n", (void*)&_executable_start);
printf("etext (end of .text): %p\n", (void*)&etext);
printf("edata (end of .data): %p\n", (void*)&edata);
printf("end (end of .bss) : %p\n", (void*)&end);
```

```
puts("\n===== High-level regions (addresses) =====");
printf("_executable_start: %p\n", (void*)&_executable_start);
printf("etext (end of .text): %p\n", (void*)&etext);
printf("edata (end of .data): %p\n", (void*)&edata);
printf("end (end of .bss): %p\n", (void*)&end);
```

```
puts("\n----- Probes (examples in those regions) -----");
```

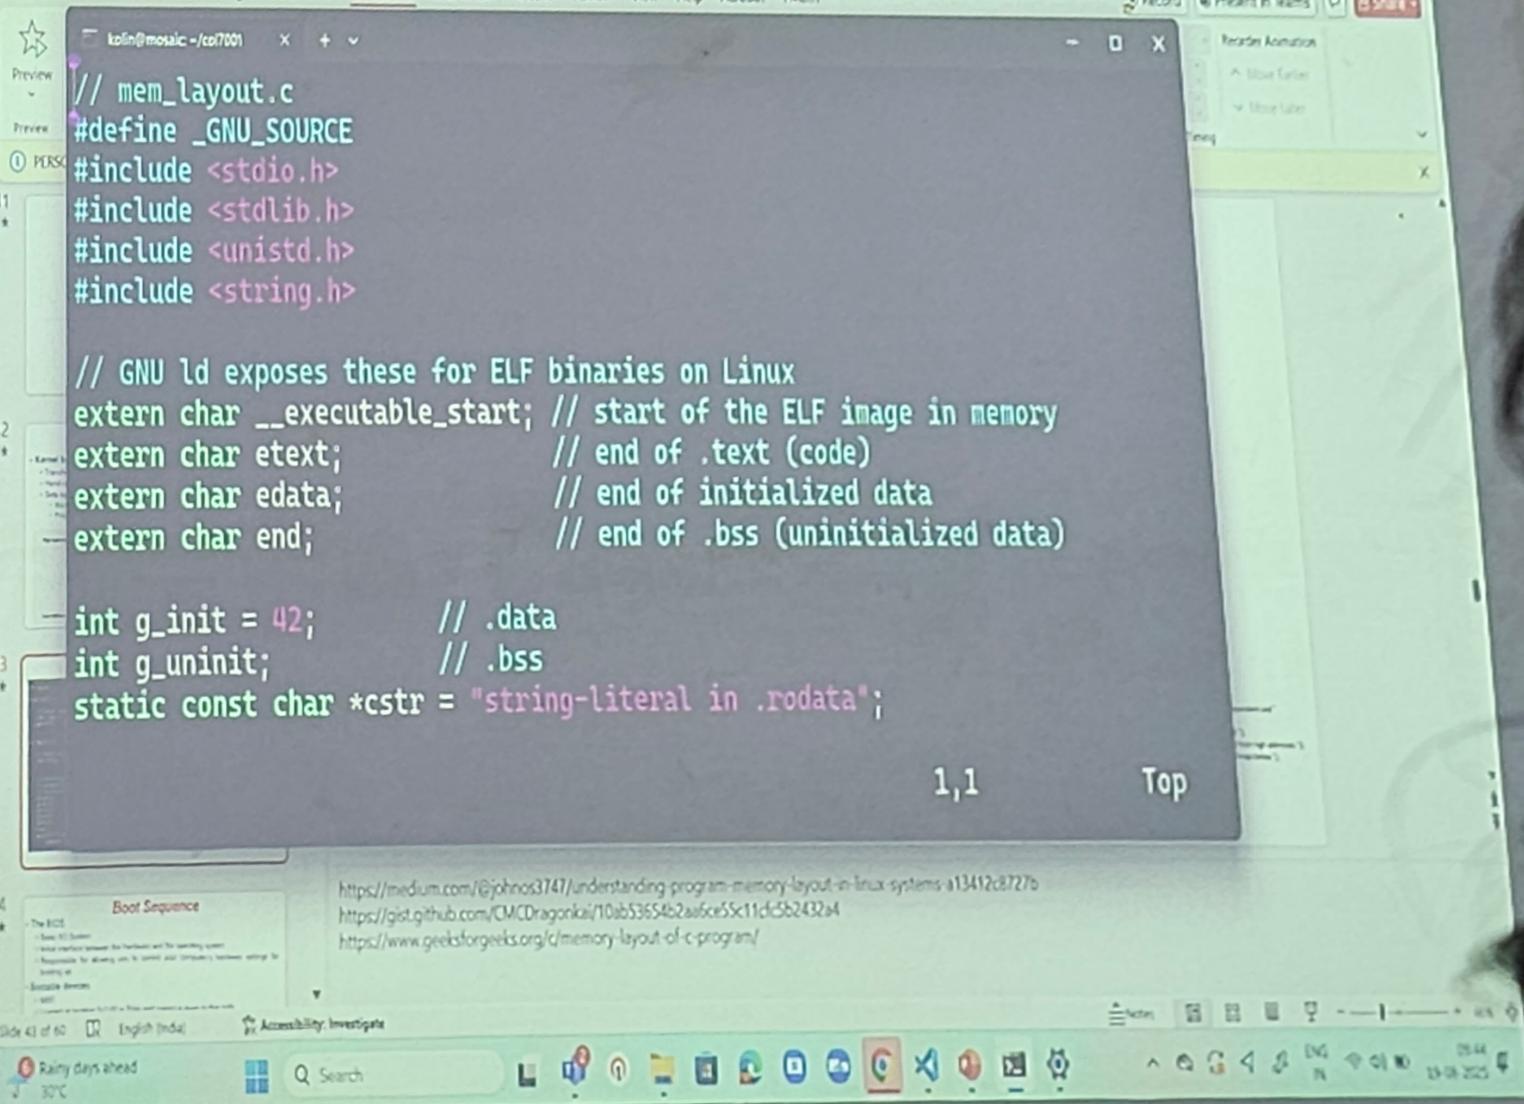
```
printf(".text main : %p\n", (void*)&main);
printf(".text foo : %p\n", (void*)&foo);
printf(".rodata cstr : %p -> \"%s\"\n", (void*)cstr, cstr);
printf(".data g_init : %p (= %d)\n", (void*)&g_init, g_init);
printf(".bss g_uninit : %p\n", (void*)&g_uninit);
printf("heap malloc(1k) : %p, %p\n", heap1, heap2);
printf("heap program brk: %p (sbrk(0))\n", brk_now);
printf("stack local var : %p (&local_on_stack)\n", (void*)&local_on_stack);
```

```
puts("\n----- Layout notes -----");
```

```
* On most x86_64 Linux builds with PIE, the executable is position-independent and  
the loader maps segments at randomized base addresses (ASLR).;
```

```
* _executable_start <= .text <= etext <= .data <= edata <= .bss <= end.;
```

```
* Heap grows upward from the program break; stack grows downward from high ad
```



Lecture 05 • Saved to this PC

File Home Insert Draw Design Transitions Animations Slide Show Record Review View Help Acrobat FMath

Record

```
kolin@mosaic:~/col7001
```

```
static const char *cstr = "string-literal in .rodata";  
static int foo(int x) { // lives in .text  
    return x + 1;  
}  
  
static void print_maps(void) {  
    puts("\n===== /proc/self/maps =====");  
    FILE *f = fopen("/proc/self/maps", "r");  
    if (!f) {  
        perror("fopen(/proc/self/maps)");  
        return;  
    }  
    char line[512];  
    while (fgets(line, sizeof line, f)) fputs(line, stdout);  
    fclose(f);  
}
```

21,0-1 25%

Boot Sequence

- The BIOS
- Boot ROM
- Initial interface between the hardware and the operating system
- Responsible for allowing the OS to control your computer's hardware settings for booting up
- Bootloader
- OS

<https://medium.com/@johnos3747/understanding-program-memory-layout-in-linux-systems-a13412c8727b>
<https://gist.github.com/CMCDragonkai/10ab53654b2aa6ce55c11cf5b2432a4>
<https://www.geeksforgeeks.org/c-memory-layout-of-c-program/>

Slide 43 of 60 English (India) Accessibility: Invert colors Notes

Rainy days ahead 30°C

Search

File Home Insert Draw Design Transitions Animations Slide Show Record Review View Help Acrobat FMath

kolin@mosaic:~/col7001

```
printf("PID: %d\n", getpid());
printf("Page size: %ld bytes\n", sysconf(_SC_PAGESIZE));

puts("\n===== High-level regions (addresses) =====");
printf("__executable_start : %p\n", (void*)&__executable_start);
printf("etext (end of .text): %p\n", (void*)&etext);
printf("edata (end of .data): %p\n", (void*)&edata);
printf("end (end of .bss) : %p\n", (void*)&end);

puts("\n----- Probes (examples in those regions) -----");
printf(".text    main      : %p\n", (void*)&main);
printf(".text    foo       : %p\n", (void*)&foo);
printf(".rodata  cstr      : %p -> \"%s\"\n", (void*)cstr, cstr);
printf(".data    g_init    : %p (= %d)\n", (void*)&g_init, g_init);
printf(".bss     g_uninit   : %p\n", (void*)&g_uninit);
printf("heap    malloc(1k) : %p, %p\n", heap1, heap2);
printf("heap    program brk: %p (sbrk(0))\n", brk_now);
```

File Home Insert Draw Design Transitions Animations Slide Show Record Review View Help Acrobat FMath

Record

Present in Teams

Share

```
tolin@mosaic:~/cat001  x + v

print_maps();

free(heap2);
free(heap1);
return 0;
}
```

Reorder Animation

▲ Move Earlier
▼ Move Later

41

42

43

44

Boot Sequence

<https://medium.com/@johnos3747/understanding-program-memory-layout-in-linux-systems-a13412c87276>

<https://gist.github.com/CMCDragonkai/10ab53654b2a1fce55x11dc5b2432a4>

<https://www.geeksforgeeks.org/c/memory-layout-of-c-program/>



Preview



Preview



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO



PERSO

Executable: /home/kolin/col7001/t2

PID: 417841

Page size: 4096 bytes

===== High-level regions (addresses) =====

__executable_start : 0x555555554000
etext (end of .text): 0x55555555570d
edata (end of .data): 0x555555558020
end (end of .bss) : 0x555555558030

----- Probes (examples in those regions) -----

.text main : 0x555555553c2
.text foo : 0x555555552e9
.rodata cstr : 0x555555556008 -> "string-literal in .rodata"
.data g_init : 0x555555558010 (=42)
.bss g_uninit : 0x55555555802c
heap malloc(1k) : 0x5555555592a0, 0x5555555596b0
:

Autosave Off Lecture-05 - Saved to this PC Search

File Home Insert Draw Design Transitions Animations Slide Show Record Review View Help Acrobat FMath

Record Present in Teams Share

Preview PERSON

41 addresses.

* For the authoritative mapping (including shared libs), see /proc/self/map s below.

```
===== /proc/self/maps =====
55555554000-55555555000 r--p 00000000 08:12 46800397
ome/kolin/col7001/t2 /h
55555555000-555555556000 r-xp 00001000 08:12 46800397
ome/kolin/col7001/t2 /h
555555556000-555555557000 r--p 00002000 08:12 46800397
ome/kolin/col7001/t2 /h
555555557000-555555558000 r--p 00002000 08:12 46800397
ome/kolin/col7001/t2 /h
555555558000-555555559000 rw-p 00003000 08:12 46800397
ome/kolin/col7001/t2 /h
555555559000-555555557a000 rw-p 00000000 00:00 0
eap] [h
```

42

43

44

Boot Sequence https://medium.com/@johns3747/understanding-program-memory-layout-in-linux-systems-a13412c8727b
https://gist.github.com/CMCDragonkai/10ab53654b2aa6ce55c11dc5b2432a4
https://www.geekforgeeks.org/c/memory-layout-of-c-program/

30°C Haze

Search

File 41 of 50 English India Accessibility Insights

Now Next Previous Last

ENG IN 08:48 19-06-2023