

What is the Course About?

- (Systems) Knowledge is a great enabler!
 - Power - 😊
 - How hardware (processors, memories, disk drives, network infrastructure) plus software (operating systems, compilers, libraries, network protocols) combine to support the execution of application programs
 - How you as a programmer can best use these resources
 - Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance

- Int and Reals
 - Not really integers and reals
 - Is $x^2 \geq 0$?

`a += b;` Actual: $a \approx 1000.220703$

- Lets see
- What about $x+y+z$
 - $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
 - $1e20 + (-1e20 + 3.14) \rightarrow$

• Computer Arithmetic

- Does not generate random values
 - Arithmetic operations have important mathematical properties
- Cannot assume all "usual" mathematical properties
 - Due to finiteness of representations
 - Integer operations satisfy "ring" properties
 - Commutativity, associativity, distributivity
 - Floating point operations satisfy "ordering" properties
 - Monotonicity, values of signs
- Observation
 - Need to understand which abstractions apply in which contexts
 - Important issues for compiler writers and serious application programm

Programmer Centric Course

- By knowing more about the underlying system, you can be more effective as a programmer
 - Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
 - Cover material in this course that you won't see elsewhere
 - Not just a course for dedicated hackers
 - We bring out the effective programmer [hidden hacker 🤖] in everyone!

Reality is More than Big-Oh

- Constant factors matter too!
- And even exact op count does not predict performance
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- Must understand system to optimize performance
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

```
void copyij(int src[2048][2048],  
            int dst[2048][2048])
```

```
{  
    for (j = 0; j < 2048; j++)  
        for (i = 0; i < 2048; i++)  
            dst[i][j] = src[i][j];  
}
```

4.3ms

```
void copyji(int src[2048][2048],  
            int dst[2048][2048])
```

```
{  
    int i, j;  
    for (j = 0; j < 2048; j++)  
        for (i = 0; i < 2048; i++)  
            dst[i][j] = src[i][j];  
}
```

81.8ms

2.0 GHz Intel Core i7 Haswell

A Program lives in Society ☹

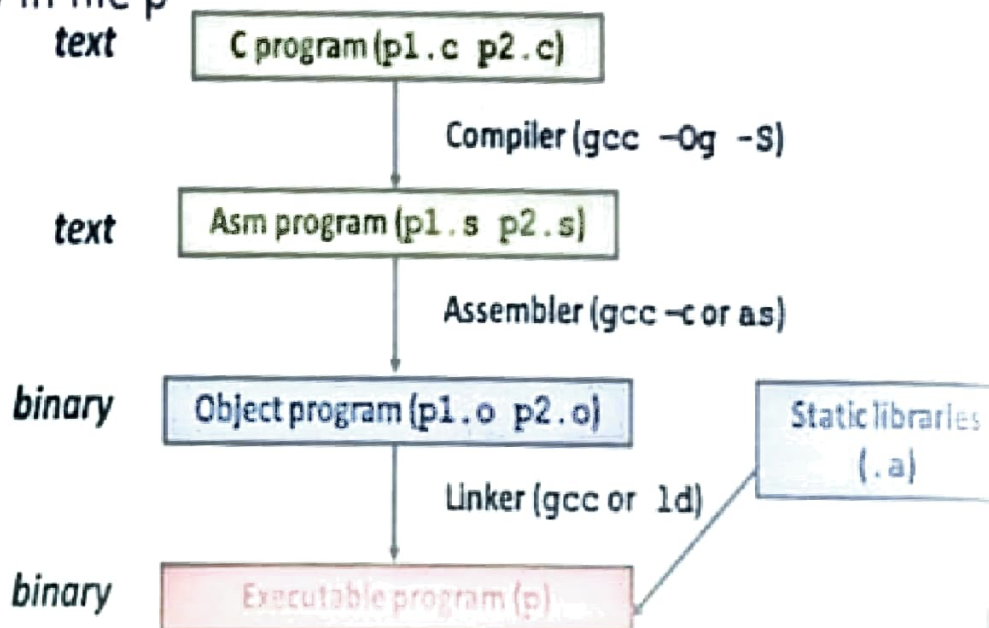
- They need to get data in and out
 - I/O system critical to program reliability and performance
- They communicate with each other over networks
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

WORDPLAY -



Turning C into Object Code

- Code in files p1.c p2.c
- Compile with command: `gcc -Og p1.c p2.c -o p`
 - Use debugging-friendly optimizations (-Og)
 - Put resulting binary in file p



Compiling Into Assembly

- Machine Specific

- gcc -Og -S compileExample.c

```
long plus(long x, long y);

void sumstore(long x, long y,
              long *dest)
{
    long t = plus(x, y);
    *dest = t;
}
```

```
file    "compileExample.c"
.text
.globl  sumstore
.type   sumstore, @function

sumstore:
.LFB0:
.cfi_startproc
endbr64
pushq   %rbx
.cfi_def_cfa_offset,16
.cfi_offset 3,-16
movq    %rdi,%rbx
call    plus@PLT
movq    %rax, (%rbx)
popq    %rbx
.cfi_def_cfa_offset,8
ret
.cfi_endproc

.LFE0:
.size   sumstore, .-sumstore
.ident  "GCC: (Ubuntu 11.4.0-1ubuntu1-23.04) 11.4.0"
.section .note.GNU-stack,"",@progbits
.section .note.gnu.property,"a"
align 8
.long  1f - 0f
.long  4f - 1f
.long  5

```

And machine Code

• Assembler

- Translates .s into .o
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

```
*dest = t;
```

```
movq %rax, (%rbx)
```

• Linker

- Resolves references between files
- Combines with static run-time libraries
 - E.g., code for **malloc**, **printf**
- Some libraries are *dynamically linked*
 - Linking occurs when program begins execution

■ C Code

- Store value *t* where designated by *dest*

■ Assembly

- Move 8-byte value to memory
 - Quad words in x86-64 parlance
- Operands:
 - t*: Register *%rax*
 - dest*: Register *%rbx*
 - **dest*: Memory *M[%rbx]*


■ Object Code

```
0x40059a: 48 89 03
```

- 3-byte instruction
- Stored at address 0x40059a

What is an OS?

- Something to do with
 - Memory Management
 - I/O Management
 - CPU Scheduling
 - Communications?
 - Does Email belong in OS?
 - Browser
 - Multitasking/multiprogramming?
- No universal definition
 - Kernel

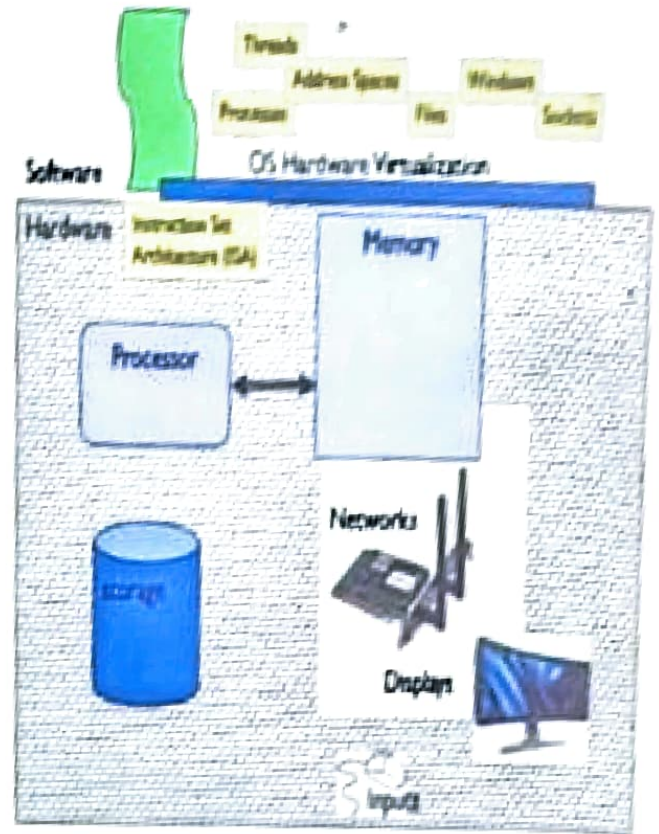



OS in a Nutshell

Computer Architecture


Magician

- Create an illusion
 - Infinite memory, dedicated machine
 - Higher level objects: files, users, messages
 - Masking limitations, virtualization




What is an OS?

- What is an Operating System?
 - And – what is it not?
- What makes Operating Systems so exciting?
- Is this a relevant course in the ML Era?
 - Actually Yes
 - Some of you will actually design and build operating systems or components of them.
 - Many of you will create systems that utilize the core concepts in operating systems.
 - All of you will build applications, etc. that utilize operating systems



Then we Need to Switch

- Context Switch
 - Referee Role
 - Manage sharing of resources, Protection, Isolation
 - Resource allocation, isolation, communication



So What is an OS?

- Layer of software that provides application software access to hardware resources
 - Abstraction the hardware
 - Provides Protection to shared resources
 - Access
 - Application
 - Security and authentication
 - Mechanisms for Communication



So How should you be Prepared

- OS is all about Programming (in C)
 - Why?
 - Efficiency
 - Wasted cycles
- Hence this require you to be very comfortable with programming and debugging C
 - Pointers (including function pointers, void*)
 - Memory Management (malloc, free, stack vs heap)
 - Debugging with GDB

```
int main(int argc, char *argv[]) {
    int i;
    for (i = 0; i < argc; i++) {
        printf("%s\n", argv[i]);
    }
    return 0;
}

int main(int argc, char *argv[]) {
    int i;
    for (i = 0; i < argc; i++) {
        printf("%s\n", argv[i]);
    }
    return 0;
}

int main(int argc, char *argv[]) {
    int i;
    for (i = 0; i < argc; i++) {
        printf("%s\n", argv[i]);
    }
    return 0;
}

int main(int argc, char *argv[]) {
    int i;
    for (i = 0; i < argc; i++) {
        printf("%s\n", argv[i]);
    }
    return 0;
}

int main(int argc, char *argv[]) {
    int i;
    for (i = 0; i < argc; i++) {
        printf("%s\n", argv[i]);
    }
    return 0;
}
```



be Prepared

```
int sum(int num1, int num2);  
int sub(int num1, int num2);  
int mult(int num1, int num2);  
int div(int num1, int num2);  
int main()
```

A

```
{ int x, y, choice, result;  
  int (*ope[4])(int, int);
```

B

```
  ope[0] = sum;  
  ope[1] = sub;  
  ope[2] = mult;  
  ope[3] = div;
```

C

```
  printf("Enter two integer numbers: ");
```

```
  scanf("%d%d", &x, &y);
```

```
  printf("Enter 0 to sum, 1 to subtract, 2 to multiply, or 3 to divide");
```

```
  scanf("%d", &choice);
```

```
  result = ope[choice](x, y);
```

```
  printf("%d", result);
```

E

```
  return 0;}
```

```
int sum(int x, int y) {return(x + y);}
```


Tying Things Together

• Glue

• Common services

- Storage, Window system, Networking
- Sharing, Authorization
- Look and feel

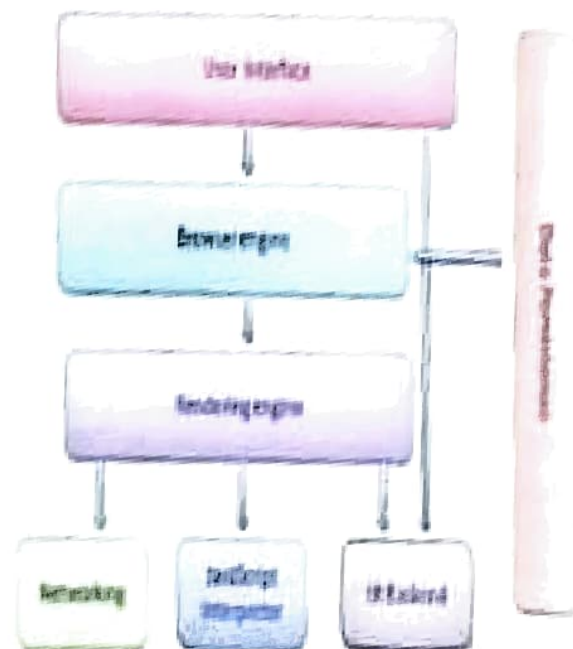



Figure: Browser components



Oops - New Terms

- Process

- Address Space
- One or more *threads* of control
- Additional system state associated with it



- Thread:

- locus of control (PC)
 - Its registers (processor state when running)
- And its "stack" (SP)
 - As required by programming language runtime

[illegible]

What is an OS?

- What is an Operating System?
 - And – what is it not?
- What makes Operating Systems so exciting?
- Is this a relevant course in the ML Era?
 - Actually Yes
 - Some of you will actually design and build operating systems or components of them.
 - Many of you will create systems that utilize the core concepts in operating systems.
 - All of you will build applications, etc. that utilize operating systems
 - The better you understand their design and implementation, the better use you'll make of them



Bare Metal to Display - 😊

- Lets write
 - test.S
- Compile
 - gcc
- Link
 - ld
- Create a bootable USB
 - dd if=/dev/zero of=usb.img bs=512 count=2880
- Copy the test.bin file to the image
 - dd if=test.bin of=usb.img
- Execute
 - Any Virtual Machine
 - Box
 - bochs
 - Homework (non Graded)
 - virtualBox or VMWare

```
.code16                                #generate 16-bit code
.text                                  #executable code location

.globl _start;
_start: #code entry point
    movb $'X', %al                    #character to print
    movb $0x0e, %ah                   #bios service code to print
    int $0x10                         #interrupt the cpu now
    . = _start + 510                   #mov to 510th byte from 0 pos
    .byte 0x55                         #append boot signature
    .byte 0xaa                         #append boot signature

as test.S -o test.o
ld -Ttext 0x7c00 --oformat=binary test.o -o test.bin
```

But Assembly – surely a Joke

- So lets try C
 - Performance
- Build the executable
 - Compile
 - Link
 - Copy

```
/*generate 16-bit code*/
__asm__(".code16\n");
/*jump boot code entry*/
__asm__("jmp $0x0000, $main\n");
/* user defined function to print series of characters terminated by null character */
void printString(const char* pStr) {
    while(*pStr) {
        __asm__ __volatile__ ("int $0x10" :: "a"(0x0e00 | *pStr), "b"(0x0007) );
        ++pStr; }
}
void main() {
    /* calling the printString function passing string as an argument */
    printString("Hello, World");
}
```

- gcc -c -g -Os -march=i686 -ffreestanding -Wall -Werror test.c -o test.o
- ld -static -Ttest.ld -nostdlib --nmagic -o test.elf test.o
- objcopy -O binary test.elf test.bin

Steps for Introducing a System Call

- Ensure that you are in the linux-xx.xx.y directory
- Create your own directory
 - mkdir info
 - cd info/
- Create processInfo.h
 - `asmlinkage long sys_listProcessInfo(void);`
- Create processInfo.c
- Write Makefile and change top level Makefile
 - `obj-y:=listProcessInfo.o`
 - `core -y += kernel/ mm/ fs/ ipc/ security/ crypto/ block/ info/`
- Alter the /arch/x86/entry/syscalls/syscall_64.tbl
 - 548 common processInfo sys_processInfo
- Alter /include/linux/syscalls.h
 - Add `asmlinkage long sys_hello(void)`
- Compile
 - `sudo make && sudo make modules_install && sudo make install`
- Test

```
#include <linux/kernel.h>
#include <linux/fs.h>
#include <linux/sched.h>
#include <linux/syscalls.h>
#include "processInfo.h"

asmlinkage long sys_listProcessInfo(void) { struct task_struct *proc;
for_each_process(proc) {
    printk(
        "Process: %d\n",
        PID_Number: %d\n",
        Process State: %d\n",
        Priority: %d\n",
        RT_Priority: %d\n",
        Static Priority: %d\n",
        Normal Priority: %d\n", );
    proc->comm;
    (long)task_pid_nr(proc);
    (long)proc->state;
    (long)proc->prio;
    (long)proc->rt_prio;
    (long)proc->static_prio;
    (long)proc->normal_prio;
    }

if(proc->parent)
    printk(
        "Parent process: %d\n",
        PID_Number: %d\n",
        proc->parent->comm;
    (long)task_pid_nr(proc->parent);
    }

    printk("n");
}
}
```

