



How to improve performance?

- There are **3** factors:
 - IPC, #instructions, and frequency
 - #instructions is dependent on the compiler → not on the architecture
- Let us look at **IPC** and **frequency**
- IPC
 - What is the IPC of an in-order pipeline?

1 if there are no stalls, otherwise < 1

Methods to
increase IPC

Forwarding

Having more not-taken branches in the code

Faster instruction and data memories



What about frequency?

- What is frequency dependent on ...
 - $\text{Frequency} = 1 / \text{clock period}$
- Clock Period:
 - 1 pipeline stage is expected to take 1 clock cycle
 - Clock period = maximum latency of the pipeline stages
- How to reduce the clock period?
 - Make each stage of the pipeline smaller by increasing the number of pipeline stages
 - Use faster transistors



Limits to Increasing Frequency

- Assume that we have the fastest possible transistors
- Can we increase the frequency to 100 GHz?

Reasons





Limits to increasing frequency - II

- What does it mean to have a very high frequency?
- Before answering, keep these facts in mind:

1

**Thumb
Rule**

$$P \propto f^3$$

 $P \rightarrow$ power $f \rightarrow$ frequency

2

**Thermo-
dynamics**

$$\Delta T \propto P$$

 $T \rightarrow$ Temperature

3

We need to increase the number of pipeline stages \rightarrow
more hazards, more forwarding paths



Pipeline Stages vs IPC

- $CPI = \frac{1}{IPC}$

$$CPI = CPI_{ideal} + stall_rate * stall_penalty$$

- The **stall rate** will remain more or less constant per instruction with the number of pipeline stages
- The **stall penalty** (in terms of cycles) will however increase
- This will lead to a net **increase** in CPI and **loss** in IPC

As we increase the number of stages, the IPC goes down.



What is ILP = Instruction level parallelism

- *multiple operations (or instructions) can be executed in parallel, from a single instruction stream*
 - *so we are not yet talking about MIMD, multiple instruction streams*

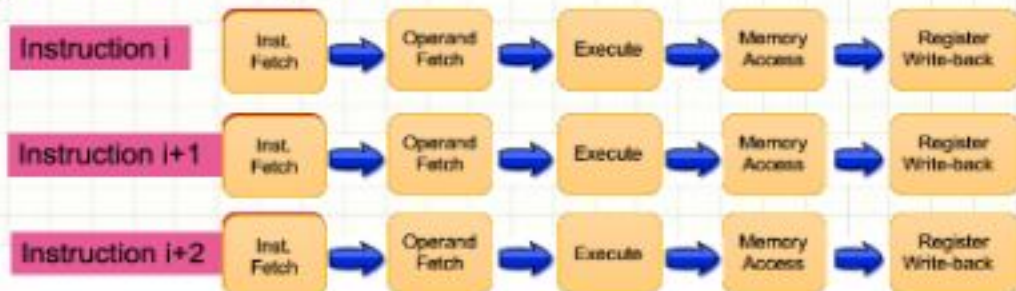
Needed:

- Sufficient (HW) resources
- Parallel scheduling
 - Hardware solution
 - Software solution
- Application should contain sufficient ILP



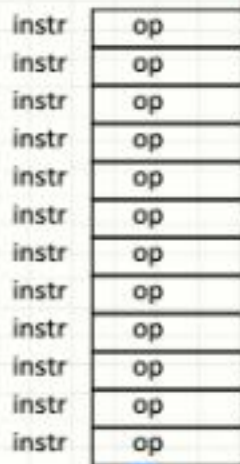
Since we cannot increase frequency ...

- Increase IPC
 - Issue **more** instructions per cycle
 - 2, 4, or 8 instructions
- Make it a **superscalar** processor → A processor that can execute multiple instructions per cycle
 - Have multiple in-order pipelines





Single Issue RISC vs Superscalar



execute
1 instr/cycle



(1-issue)
RISC CPU

Change HW,

but can use
same code



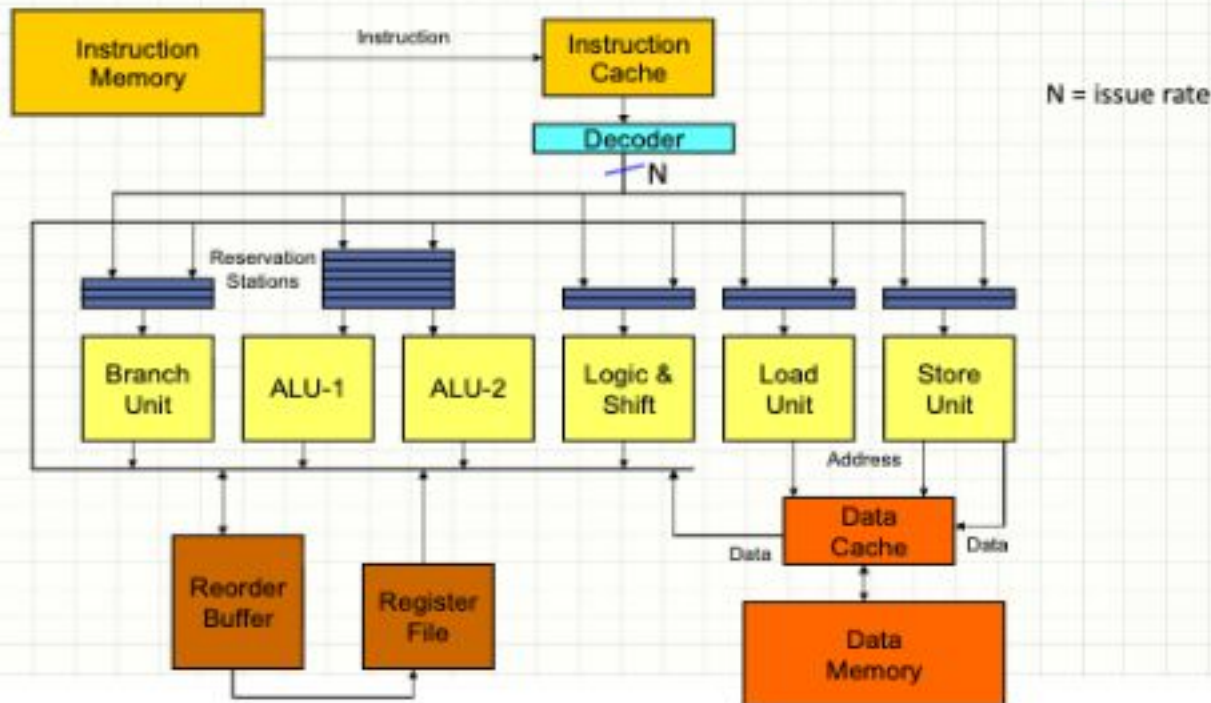
3-issue Superscalar

issue and (try to) execute
3 instr/cycle





Superscalar: General Architecture Concept





Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX				
SUB.D	F8, F2, F6		IF	EX				
DIV.D	F10, F0, F6			IF				
ADD.D	F6, F8, F2			IF				
MUL.D	F12, F2, F4							



Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX	EX			
SUB.D	F8, F2, F6		IF	EX	EX			
DIV.D	F10, F0, F6			IF				
ADD.D	F6, F8, F2			IF				
MUL.D	F12, F2, F4							

stall because
of data dep.

cannot be fetched because window full



Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX	EX	EX		
SUB.D	F8, F2, F6		IF	EX	EX	WB		
DIV.D	F10, F0, F6			IF				
ADD.D	F6, F8, F2			IF		EX		
MUL.D	F12, F2, F4					IF		



Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX	EX	EX	EX	
SUB.D	F8, F2, F6		IF	EX	EX	WB		
DIV.D	F10, F0, F6			IF				
ADD.D	F6, F8, F2			IF		EX	EX	
MUL.D	F12, F2, F4					IF		

cannot execute
structural hazard



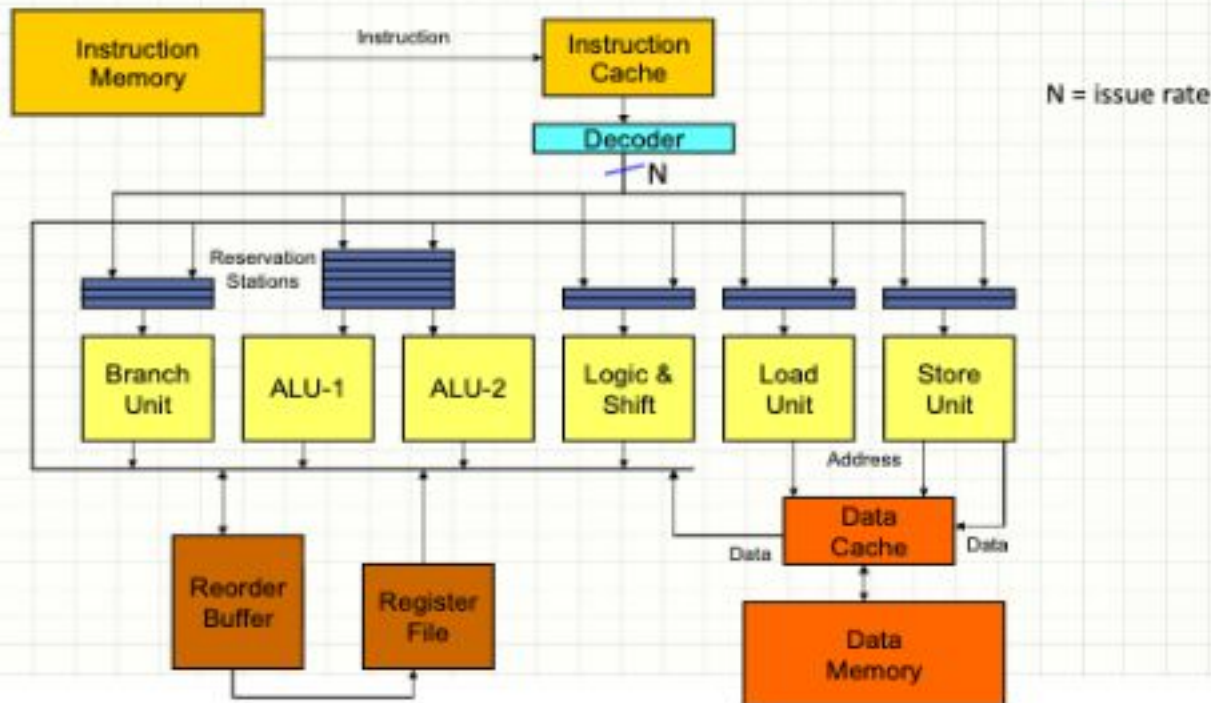
Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX	EX	EX	EX	WB
SUB.D	F8, F2, F6		IF	EX	EX	WB		
DIV.D	F10, F0, F6			IF				EX
ADD.D	F6, F8, F2			IF		EX	EX	WB
MUL.D	F12, F2, F4					IF		?



Superscalar: General Architecture Concept





Hazards

- Three types of hazards
 - **Structural**
 - Multiple instructions need access to the same hardware at the same time
 - **Data dependence**
 - There is a dependence between operands (in register or memory) of successive instructions
 - **Control dependence**
 - Determines the order of the execution of basic blocks
 - When jumping/branching to another address the pipeline has to be (partly) squashed and refilled
- Hazards cause scheduling problems and delay the pipeline



Data dependences

- **RaW** read after write
 - real or flow dependence
 - can only be avoided by value prediction (i.e. speculating on the outcome of a previous operation)
- **WaR** write after read
- **WaW** write after write
 - WaR and WaW are **false or name dependencies**
 - Could be avoided by **renaming** (if sufficient registers are available); see later slide

Notes:

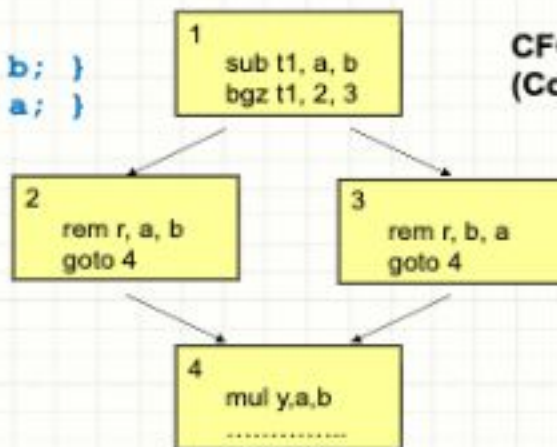
1. data dependences can be **both** between **registers** and **memory** data operations
2. data dependencies are shown in the DDG: Data Dependence Graph



Control Dependences: CFG

C input code:

```
if (a > b) { r = a % b; }  
else      { r = b % a; }  
y = a*b;
```



Questions:

- How real are control dependences?
- Can 'mul y, a, b' be moved to block 2, 3 or even block 1?
- Can 'rem r, a, b' be moved to block 1 and executed speculatively?



Avoiding pipeline stalls due to Hazards

- **Structural**

- Buy more hardware
 - Extra units, pipelined units, more ports on RF and data memory (or banked memories), etc.
- Note: more HW means bigger chip => could increase cycle time t_{cycle}

- **Data dependence**

- Real (RaW) dependences: add **Forwarding** (aka **Bypassing**) logic
 - Compiler optimizations
- False (WaR & WaW) dependences: use renaming (either in HW or in SW)

- **Control dependence**

- Adding extra pipeline HW to reduce the number of Branch delay slots
- Branch prediction
- Avoiding Branches



FP Loop: Where are the Hazards?

```
for (i=1000; i>0; i=i-1)
```

```
  x[i] = x[i] + s;
```

Instruction producing result	Instruction using result	Latency in cycles	stalls between in cycles
FP ALU op	Another FP ALU op	4	3
FP ALU op	Store double	3	2
Load double	FP ALU op	1	1
Load double	Store double	1	0
Integer op	Integer op	1	0

```
Loop: L.D    F0,0(R1);F0=vector element
      ADD.D  F4,F0,F2;add scalar from F2
      S.D    0(R1),F4;store result
      DADDUI R1,R1,-8;decrement pointer 8B
      BNEZ   R1,loop;branch R1!=zero
```

```
1 Loop: L.D    F0,0(R1)
2      DADDUI R1,R1,-8
3      ADD.D  F4,F0,F2
4      stall
5      stall
6      S.D    0(R1),F4
7      BNEZ   R1,Loop
```

```
1 Loop:L.D    F0,0(R1)
2      L.D    F6,-8(R1)
3      L.D    F10,-16(R1)
4      L.D    F14,-24(R1)
5      ADD.D  F4,F0,F2
6      ADD.D  F8,F6,F2
7      ADD.D  F12,F10,F2
8      ADD.D  F16,F14,F2
9      S.D    0(R1),F4
10     S.D    -8(R1),F8
11     S.D    -16(R1),F12
12     DADDUI R1,R1,#32
13     S.D    0(R1),F16;0-32=-24
14     BNEZ   R1,LOOP
```

14 clock cycles, or 3.5 per iteration

```
1 Loop: L.D    F0,0(R1);F0=vector element
2      stall
3      ADD.D  F4,F0,F2;add scalar in F2
4      stall
5      stall
6      S.D    0(R1),F4;store result
7      DADDUI R1,R1,-8;decrement pointer 8B (DW)
8      stall;assumes can't forward to branch
9      BNEZ   R1,Loop;branch R1!=zero
```

```
1 Loop:L.D    F0,0(R1)
3      ADD.D  F4,F0,F2
6      S.D    0(R1),F4;drop DSUBUI & BNEZ
7      L.D    F6,-8(R1)
9      ADD.D  F8,F6,F2
12     S.D    -8(R1),F8;drop DSUBUI & BNEZ
13     L.D    F10,-16(R1)
15     ADD.D  F12,F10,F2
18     S.D    -16(R1),F12;drop DSUBUI & BNEZ
19     L.D    F14,-24(R1)
21     ADD.D  F16,F14,F2
24     S.D    -24(R1),F16
25     DADDUI R1,R1,#-32;alter to 4*8
26     BNEZ   R1,LOOP
```

27 clock cycles, or 6.75 per iteration

(Assumes R1 is multiple of 4)



Dynamic Scheduling Principle

- What we examined so far is *static scheduling*
 - Compiler reorders instructions so as to avoid hazards and reduce stalls
- Dynamic scheduling*:
hardware rearranges instruction execution to reduce stalls
- Example:

```
DIV.D  F0,F2,F4      ; takes 24 cycles and  
RaW: real dependence ; is not pipelined  
ADD.D  F10,F0,F8  
  
SUB.D  F12,F8,F14
```

*This instruction cannot continue
even though it does not depend
on previous Div and Add*

- Key idea:** Allow instructions behind stall to proceed



Register Renaming: General Idea

- Example, look at F6:



- False (aka “name”) dependencies with F6
 - anti: WaR and
 - output: WaW in this example

Question: how can this code (optimally) be executed?



Register Renaming Technique

- Eliminate name/false dependencies:
 - anti- (**WaR**) and output (**WaW**) dependencies
- Can be implemented
 - by the **compiler**
 - advantage: low cost
 - disadvantage: “old” codes perform poorly
 - by **hardware**
 - advantage: binary compatibility
 - disadvantage: extra hardware needed



Register Renaming by HW

- Same example after renaming:

DIV.D R, F2, F4
ADD.D S, R, F8
S.D S, 0(R1)
SUB.D T, F10, F14
MUL.D U, F10, T

Original code:

DIV.D F0, F2, F4
ADD.D F6, F0, F8
S.D F6, 0(R1)
SUB.D F8, F10, F14
MUL.D F6, F10, F8

- Each destination gets a new (physical) register assigned
- Now only RaW hazards remain, which can be strictly ordered
- We will see several HW implementations of Register Renaming
 1. use ReOrder Buffer (ROB) & Reservation Stations, or
 2. use large register file with mapping table



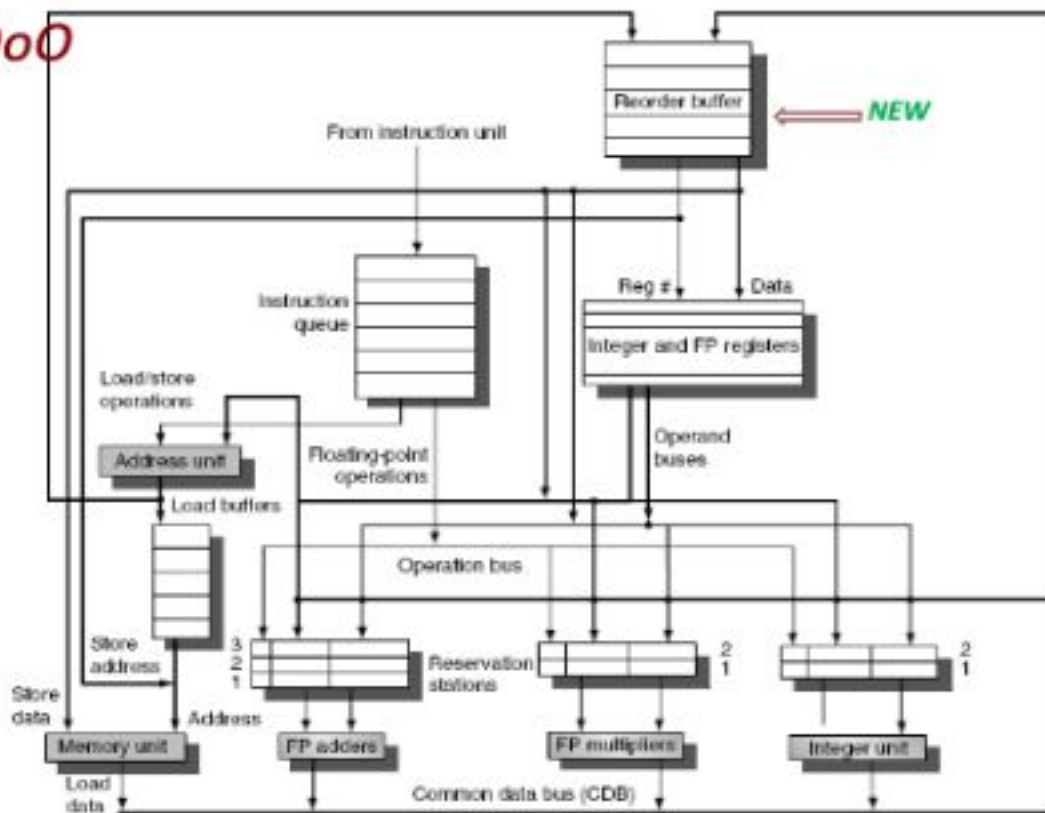
Speculation (Hardware based)

- Execute instructions along predicted execution paths **but *only commit the results if the prediction was correct***
- **Instruction commit:** *allowing an instruction to only **update** the register file when instruction is no longer speculative*
- **Need an additional piece of hardware to prevent any irrevocable action until an instruction commits:**
 - **Reorder buffer, or Large renaming register file**
 - **why? think about it?**



Speculative OoO

execution with
speculation
using RoB





Reorder Buffer (RoB)

- Register values and memory values are not written until an instruction commits
 - every destination gets a new entry in the RoB
- On misprediction:
 - Speculated entries in RoB are cleared
- Exceptions:
 - Not recognized/taken until it is ready to commit
 - Precise exceptions require that 'later' entries in RoB are cleared



Flynn* Taxonomy, 1966

- In 2013, SIMD and MIMD most common parallelism in architectures – usually both in same system!
- Most common parallel processing programming style: Single Program Multiple Data (“SPMD”)
 - Single program that runs on all processors of a MIMD
 - Cross-processor execution coordination using synchronization primitives
- SIMD (aka hw-level *data parallelism*): specialized function units, for handling lock-step calculations involving arrays
 - Scientific computing, signal processing, multimedia (audio/video processing)

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345 (Clovertown)

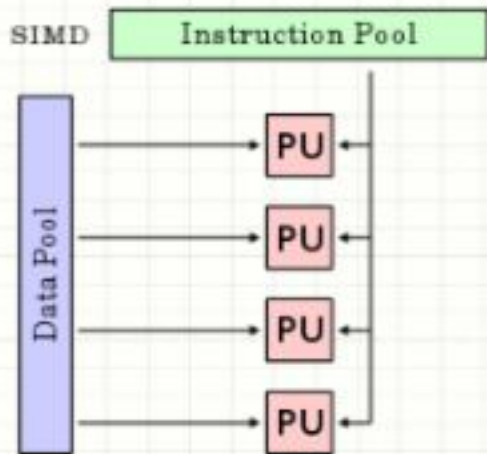
*Prof. Michael Flynn
Stanford





Single-Instruction/Multiple-Data Stream (SIMD or “sim-dee”)

- SIMD computer exploits multiple data streams against a single instruction stream to operations that may be naturally parallelized, e.g., Intel SIMD instruction extensions or NVIDIA Graphics Processing Unit (GPU)

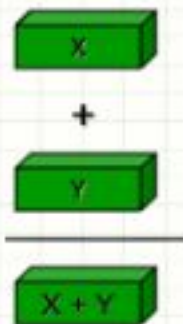




SIMD: Single Instruction, Multiple Data

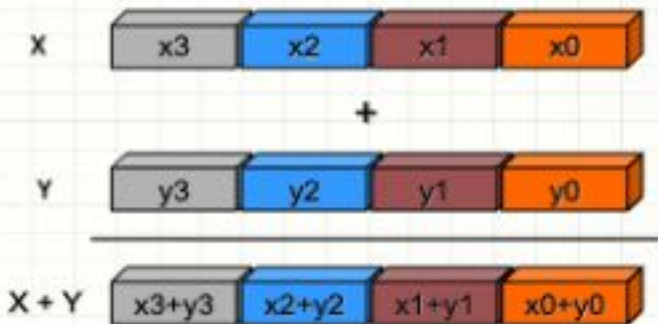
- Scalar processing

- traditional mode
- one operation produces one result



- SIMD processing

- With Intel SSE / SSE2
- SSE = streaming SIMD extensions
- one operation produces multiple results



Slide Source: Alex Klimovitski & Dean Macri, Intel Corporation



What does this mean to you?

- In addition to SIMD extensions, the processor may have other special instructions
 - Fused Multiply-Add (FMA) instructions:
$$x = y + c * z$$

is so common some processor execute the multiply/add as a single instruction, at the same rate (bandwidth) as + or * alone
- In theory, the compiler understands all of this
 - When compiling, it will rearrange instructions to get a good “schedule” that maximizes pipelining, uses FMAs and SIMD
 - It works with the mix of instructions inside an inner loop or other block of code
- But in practice the compiler may need your help
 - Choose a different compiler, optimization flags, etc.
 - Rearrange your code to make things more obvious
 - Using special functions (“intrinsics”) or write in assembly ☹



Intel SIMD Extensions

- MMX 64-bit registers, reusing floating-point registers [1992]
- SSE2/3/4, new 8 128-bit registers [1999]

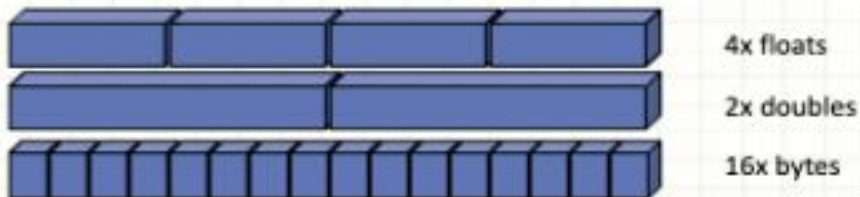
127		0
	XMM7	
	XMM6	
	XMM5	
	XMM4	
	XMM3	
	XMM2	
	XMM1	
	XMM0	

- AVX, new 256-bit registers [2011]
 - Space for expansion to 1024-bit registers

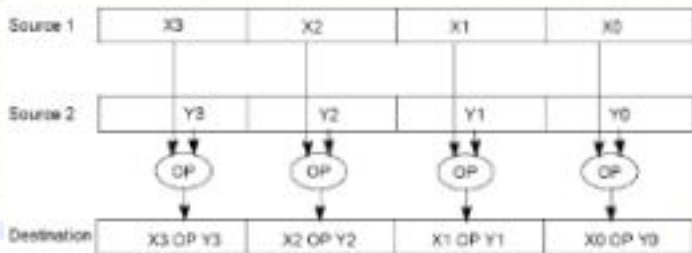


SSE / SSE2 SIMD on Intel

- SSE2 data types: anything that fits into 16 bytes, e.g.,



- Instructions perform add, multiply etc. on all the data in parallel



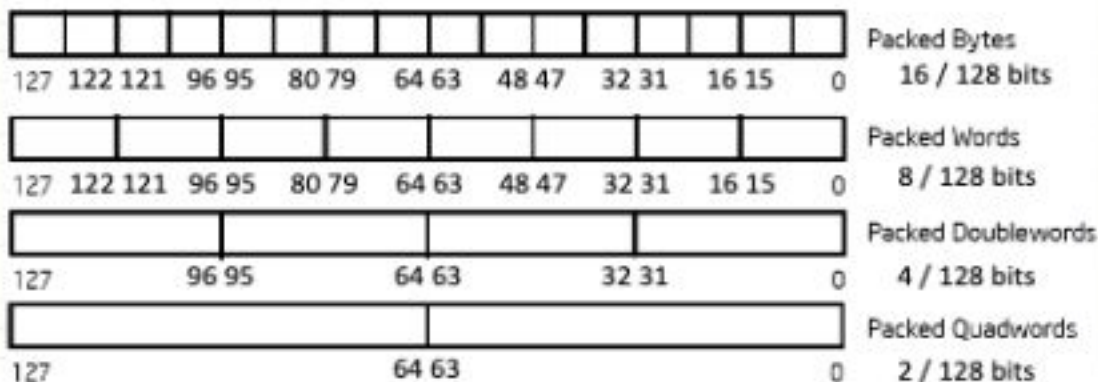
- Similar on GPUs, vector p (operations)



Intel Architecture SSE2+ 128-Bit SIMD Data Types

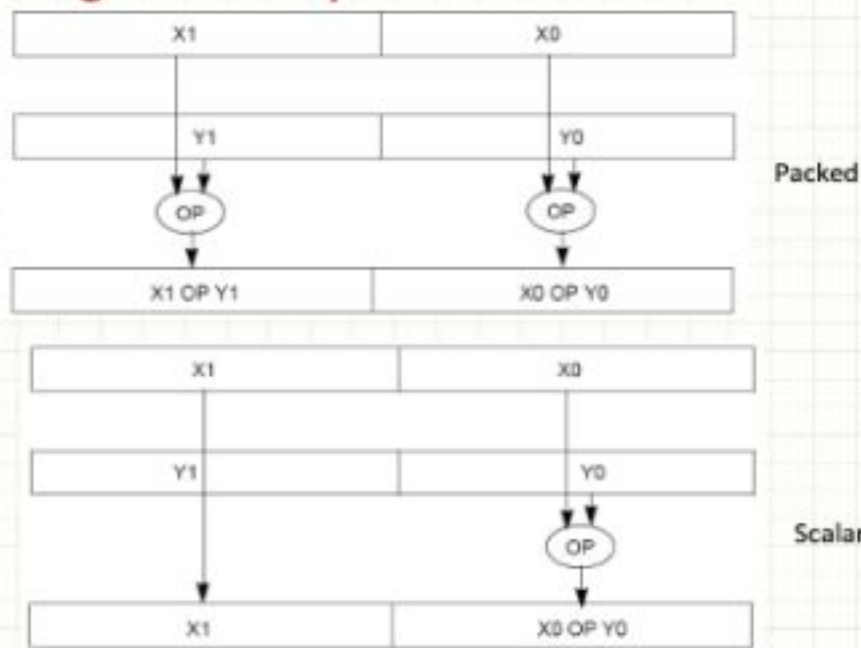
- Note: in Intel Architecture (unlike MIPS) a word is 16 bits
 - Single-precision FP: Double word (32 bits)
 - Double-precision FP: Quad word (64 bits)

Fundamental 128-Bit Packed SIMD Data Types





Packed and Scalar Double-Precision Floating-Point Operations





Example: SIMD Array Processing

```
for each f in array  
    f = sqrt(f)
```

```
for each f in array  
{  
    load f to floating-point register  
    calculate the square root  
    write the result from the  
    register to memory  
}
```

```
{  
    for each 4 members in array  
    {  
        load 4 members to the SSE register  
        calculate 4 square roots in one operation  
        store the 4 results from the register to memory  
    }  
}
```

SIMD style



Data-Level Parallelism and SIMD

- SIMD wants adjacent values in memory that can be operated in parallel
- Usually specified in programs as loops

```
for(i=1000; i>0; i=i-1)  
    x[i] = x[i] + s;
```

- How can reveal more data-level parallelism than available in a single iteration of a loop?
- *Unroll loop* and adjust iteration rate



Loop Unrolling in C

- Instead of compiler doing loop unrolling, could do it yourself in C

```
for (i=1000; i>0; i=i-1)
    x[i] = x[i] + s;
```

- Could be rewritten

```
for (i=1000; i>0; i=i-4) {
    x[i]      = x[i] + s;
    x[i-1]    = x[i-1] + s;
    x[i-2]    = x[i-2] + s;
    x[i-3]    = x[i-3] + s;
}
```




Building a Python-C extension

- Write the C function
 - PyObject
 - PyArg_ParseTuple()
 - PyLong_FromLong()
- Write the init function
 - PyMethodDef
 - PyModuleDef
- PyMODINIT_FUNC
 - First import calls this

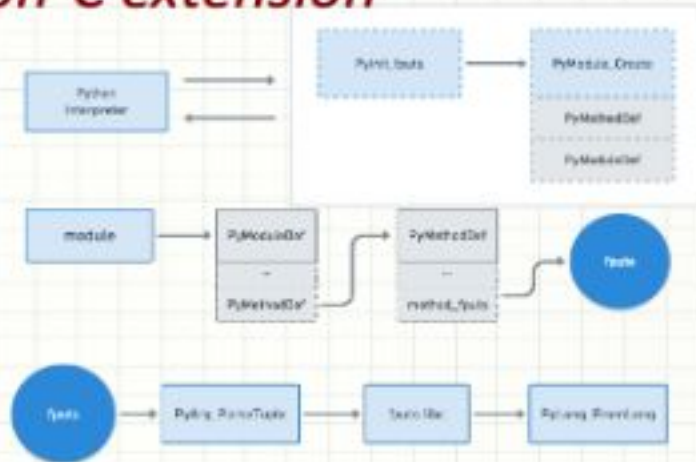
```
PyMODINIT_FUNC PyInit_fputs(void) {  
    return PyModule_Create(&fputsmodule);  
}
```

```
#include <Python.h>  
  
static PyObject *method_fputs(PyObject *self, PyObject *args) {  
    char *str, *filename = NULL;  
    int bytes_copied = -1;  
  
    /* Parse arguments */  
    if(!PyArg_ParseTuple(args, "ss", &str, &filename)) {  
        return NULL;  
    }  
  
    FILE *fp = fopen(filename, "w");  
    bytes_copied = fputs(str, fp);  
    fclose(fp);  
  
    return PyLong_FromLong(bytes_copied);  
}  
  
static PyMethodDef fputsMethods[] = {  
    {"fputs", method_fputs, METH_VARARGS, "Python interface for fputs"},  
    {NULL, NULL, 0, NULL}  
};  
  
static struct PyModuleDef fputsmodule = {  
    PyModuleDef_HEAD_INIT,  
    "fputs",  
    "Python interface for the fputs C library function",  
    -1,  
    fputsMethods  
};
```



Building a Python-C extension

- This is how it appears
- Create setup.py
- Compile and Install
 - python3_setup.py install
- Use in your code



```

>>> import fputs
>>> fputs.__doc__
'Python interface for the fputs C library function'
>>> fputs.__name__
'fputs'
>>> # Write to an empty file named 'write.txt'
>>> fputs.fputs("Real Python!", "write.txt")
13
>>> with open("write.txt", "r") as f:
>>>     print(f.read())
'Real Python!'
  
```

```
from distutils.core import setup, Extension
```

```

def main():
    setup(name="fputs",
          version="1.0.0",
          description="Python interface for the fputs C library function",
          author="your name",
          author_email="your_name@yourmail.com",
          ext_modules=[Extension("fputs", ["fputsmodule.c"])]

if __name__ == "__main__":
    main()
  
```



Accelerate matrixMultiply

```
#include <emmintrin.h>
#include <sys/time.h>
#include <stdio.h>
#include <stdint.h>
#include <stdlib.h>
#include <math.h>
#include <malloc.h>
```

```
[[
unsigned long long get_timestamp ()
{
    struct timeval now;
    gettimeofday (&now, NULL);
    return now.tv_usec + (unsigned long long)now.tv_sec * 1000000;
}
]]
```

```
void dgemv_basic(const uint32_t n, const double* A, const double* B, double* C)
{
    for(uint32_t i = 0; i < n; ++i)
    {
        for(uint32_t j = 0; j < n; ++j)
        {
            double xij = C[i + j * n]; // xij = C[i][j]
            for(uint32_t k = 0; k < n; k++)
            {
                xij += A[i + k * n] * B[k + j * n]; // xij += A[i][k]*B[k][j]
            }
            C[i + j * n] = xij; // C[i][j] = xij
        }
    }
}
```

```
void dgemv_ava256(const uint32_t n, const double* A, const double* B, double* C)
{
    for( uint32_t i = 0; i < n; i += 4 )
    {
        for( uint32_t j = 0; j < n; j++)
        {
            __m256d c0 = _mm256_load_pd(C + i + j * n); // c0 = C[i][j]
            for( uint32_t k = 0; k < n; k++)
            {
                c0 = _mm256_add_pd(c0, A[i][k]*B[k][j]);
                _mm256_store_pd( _mm256_load_pd(C + i + k * n), _mm256_broadcast_sd(B + k + j * n) );
            }
            _mm256_store_pd(C + i + j * n, c0); // C[i][j] = c0
        }
    }
}
```

```
void dgemv_ava512(const uint32_t n, const double* A, const double* B, double* C)
{
    for( uint32_t i = 0; i < n; i += 8 )
    {
        for( uint32_t j = 0; j < n; ++j)
        {
            __m512d c0 = _mm512_load_pd(C + i + j * n); // c0 = C[i][j]
            for( uint32_t k = 0; k < n; k++)
            {
                // c0 += A[i][k] * B[k][j]
                __m512d bb = _mm512_broadcast_sd( _mm_load_sd(B + k + j * n + k));
                c0 = _mm512_fmadd_pd(_mm512_load_pd(A + i + n * k + i), bb, c0);
            }
            _mm512_store_pd(C + i + j * n, c0); // C[i][j] = c0
        }
    }
}
```



Accelerate matrixMultiply

```

int main(){
    uint32_t trial_no = 11;
    uint32_t n = 32 * 20;

    double *a;
    double *b;
    double *c;
    unsigned long long t0;
    unsigned long long t1;
    unsigned long long t;

    struct timeval tv1, tv2;

    double result;

    a = (double*) _mm_malloc (n * n * sizeof(double), 64);
    b = (double*) _mm_malloc (n * n * sizeof(double), 64);
    c = (double*) _mm_malloc (n * n * sizeof(double), 64);

    for(uint32_t i = 0; i < n * n; ++i)
    {
        a[i] = rand();
    }

    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_essetm, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }
    result = result / (double)(trial_no);
    printf("Time Taken: Ref:\n", result);

    result=0;
    for(uint3_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemv_essetm, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }
    result = result / (double)(trial_no);
    printf("Time Taken: AVX256:\n", result);

    result=0;
    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemv_essetm, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }
    result = result / (double)(trial_no);
    printf("Time Taken: AVX32:\n", result);
}

```

```

kallagrinerva:~/accel/extension/tnw$ ./a.out
Time Taken: Ref:      1.027594
Time Taken: AVX256:   0.436642
Time Taken: AVX32:    0.412659

```

```
gcc -Wall -pedantic matmul-3.c -march=icelake-server
```



Using shared Objects

```
from ctypes import CDLL, POINTER  
from ctypes import c_size_t, c_double, c_uint  
import numpy as np  
from numpy import random  
import time
```

```
gcc -Wall -pedantic -shared -fPIC -o mylib.so try.c -march=icelake-server
```

```
# load the library  
mylib = CDLL("/home/kalin/aceel/extension/two/mylib.so")  
  
# C-type corresponding to numpy array  
ND_POINTER_1 = np.ctypeslib.ndpointer(dtype=np.float64,  
                                       ndim=1,  
                                       flags="C")  
  
# define prototypes  
mylib.print_array.argtypes = [ND_POINTER_1, c_size_t]  
mylib.print_array.restype = None  
mylib.dgemm_basic.argtypes = [c_uint, ND_POINTER_1, ND_POINTER_1, ND_POINTER_1]  
mylib.dgemm_basic.restype = None  
mylib.dgemm_avx512.argtypes = [c_uint, ND_POINTER_1, ND_POINTER_1, ND_POINTER_1]  
mylib.dgemm_avx512.restype = None  
mylib.dgemm_avx256.argtypes = [c_uint, ND_POINTER_1, ND_POINTER_1, ND_POINTER_1]  
mylib.dgemm_avx256.restype = None
```

```
n=32*20
```

```
#s = (2,2)  
s=n*n  
a = random.rand(s)  
b =random.rand(s)  
c =np.zeros(s)  
X=np.ones(5)
```

```
# call function  
mylib.print_array(X, X.size)  
start=time.time()  
mylib.dgemm_basic(n,a,b,c)  
end=time.time()  
print(c)  
print("Elapsed Time Ref:", end-start)  
c =np.zeros(s)  
start1=time.time()  
mylib.dgemm_avx256(n,a,b,c)  
end1=time.time()  
print("Elapsed Time AVX256:", end1-start1)
```